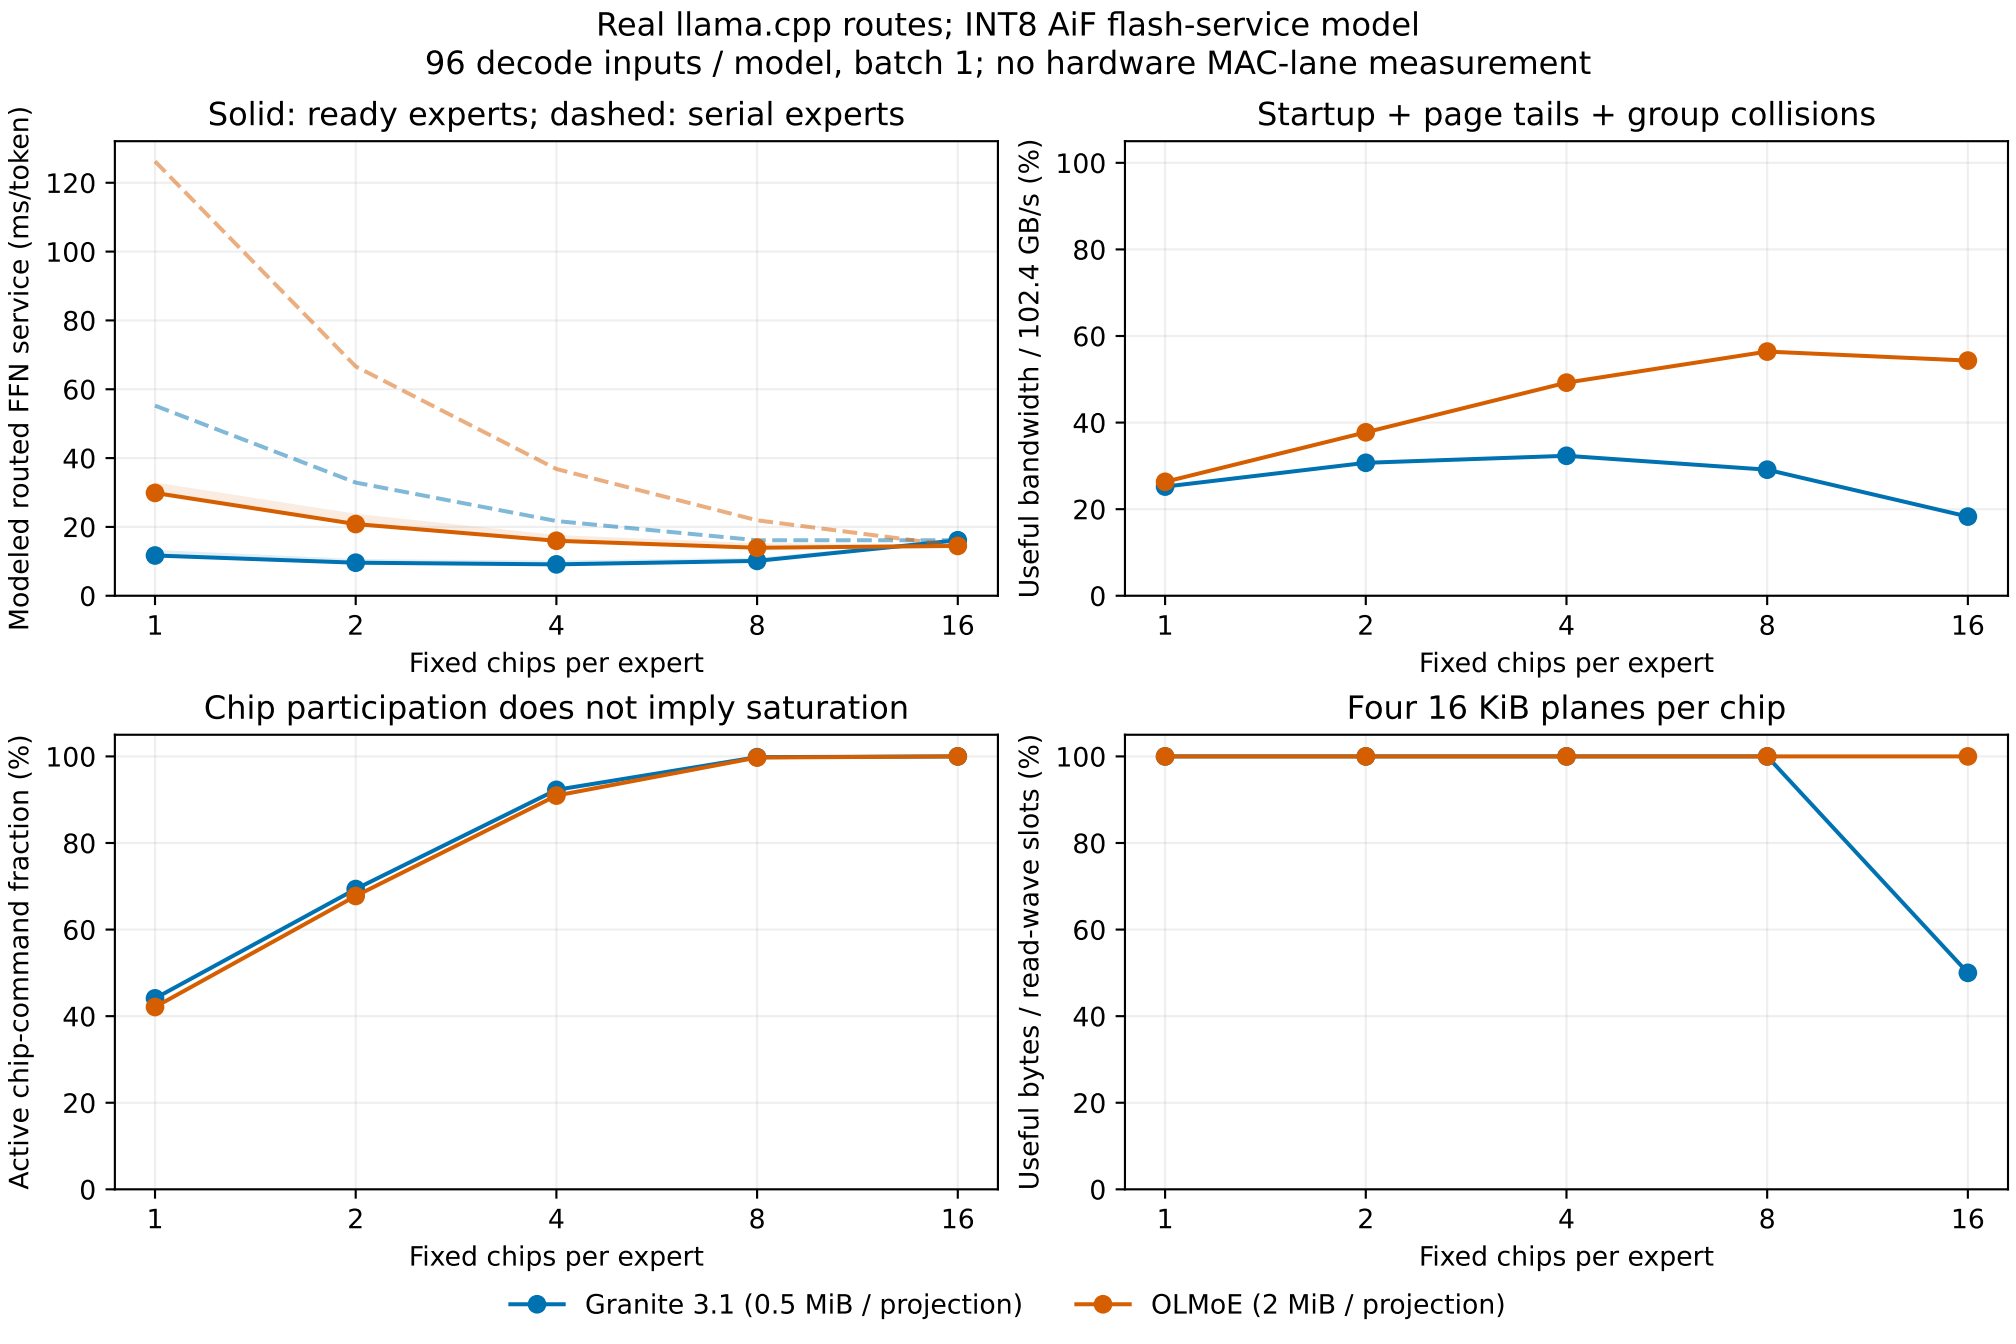
Real llama.cpp routes; INT8 AiF flash-service model
96 decode inputs / model, batch 1; no hardware MAC-lane measurement
Solid: ready experts; dashed: serial experts
Modeled routed FFN service (ms/token)
Fixed chips per expert
Startup + page tails + group collisions
Useful bandwidth / 102.4 GB/s (%)
Fixed chips per expert
Chip participation does not imply saturation
Active chip-command fraction (%)
Fixed chips per expert
Four 16 KiB planes per chip
Useful bytes / read-wave slots (%)
Fixed chips per expert
Granite 3.1 (0.5 MiB / projection)
OLMoE (2 MiB / projection)