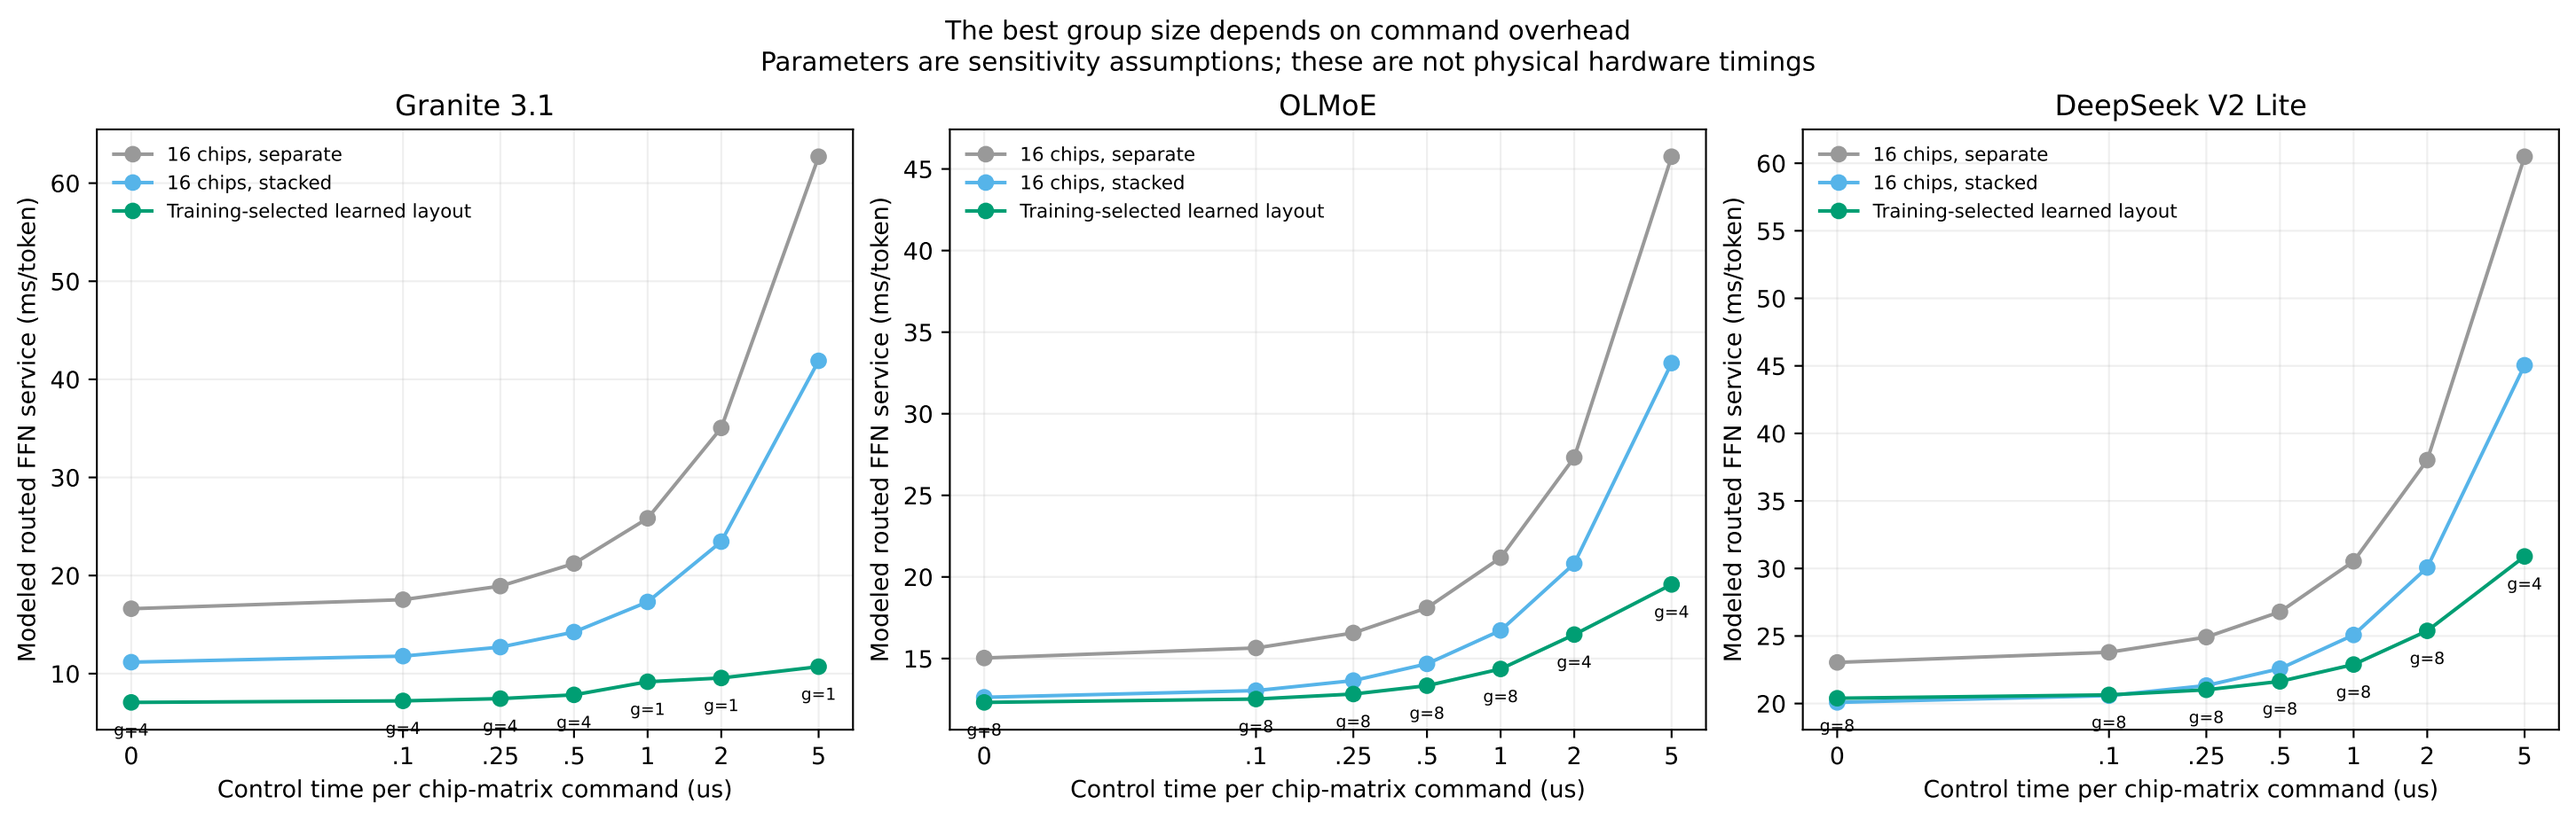
The best group size depends on command overhead
Parameters are sensitivity assumptions; these are not physical hardware timings
Granite 3.1
OLMoE
DeepSeek V2 Lite
16 chips, separate
16 chips, stacked
Training-selected learned layout
Modeled routed FFN service (ms/token)
Control time per chip-matrix command (us)
0
.1
.25
.5
1
2
5
g=4
g=1
g=8