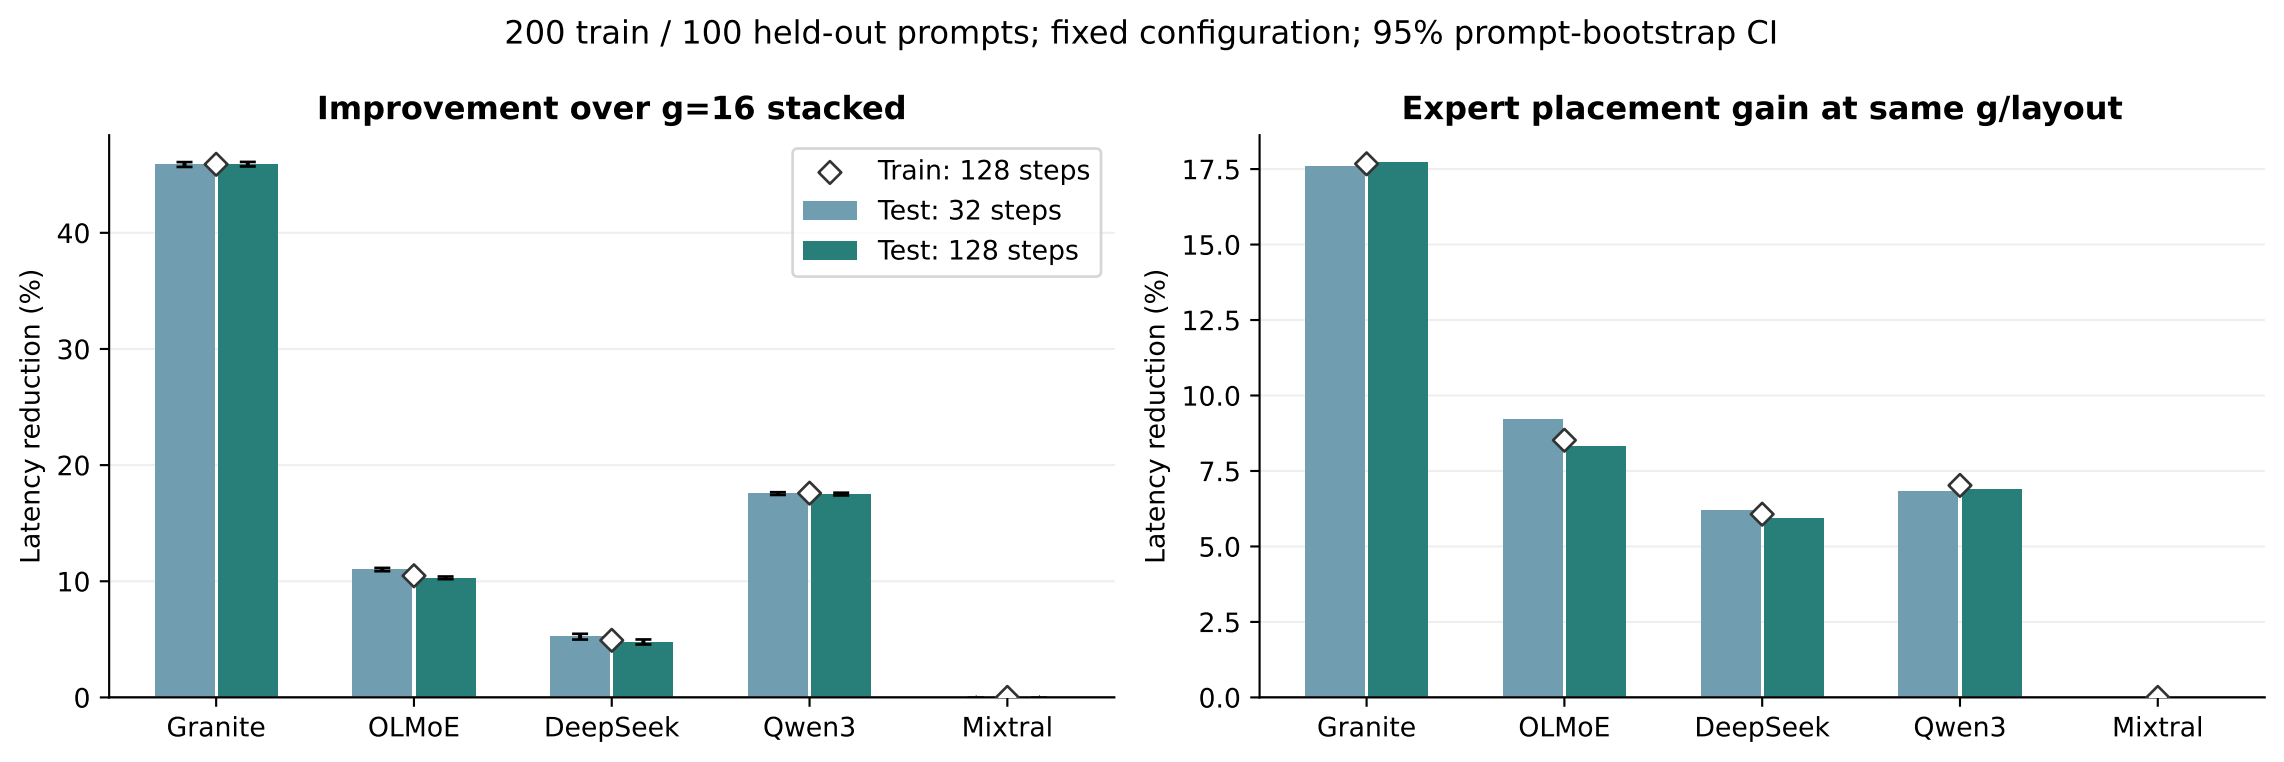
200 train / 100 held-out prompts; fixed configuration; 95% prompt-bootstrap CI
Improvement over g=16 stacked
Latency reduction (%)
0
10
20
30
40
Train: 128 steps
Test: 32 steps
Test: 128 steps
Granite
OLMoE
DeepSeek
Qwen3
Mixtral
Expert placement gain at same g/layout
Latency reduction (%)
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
Granite
OLMoE
DeepSeek
Qwen3
Mixtral