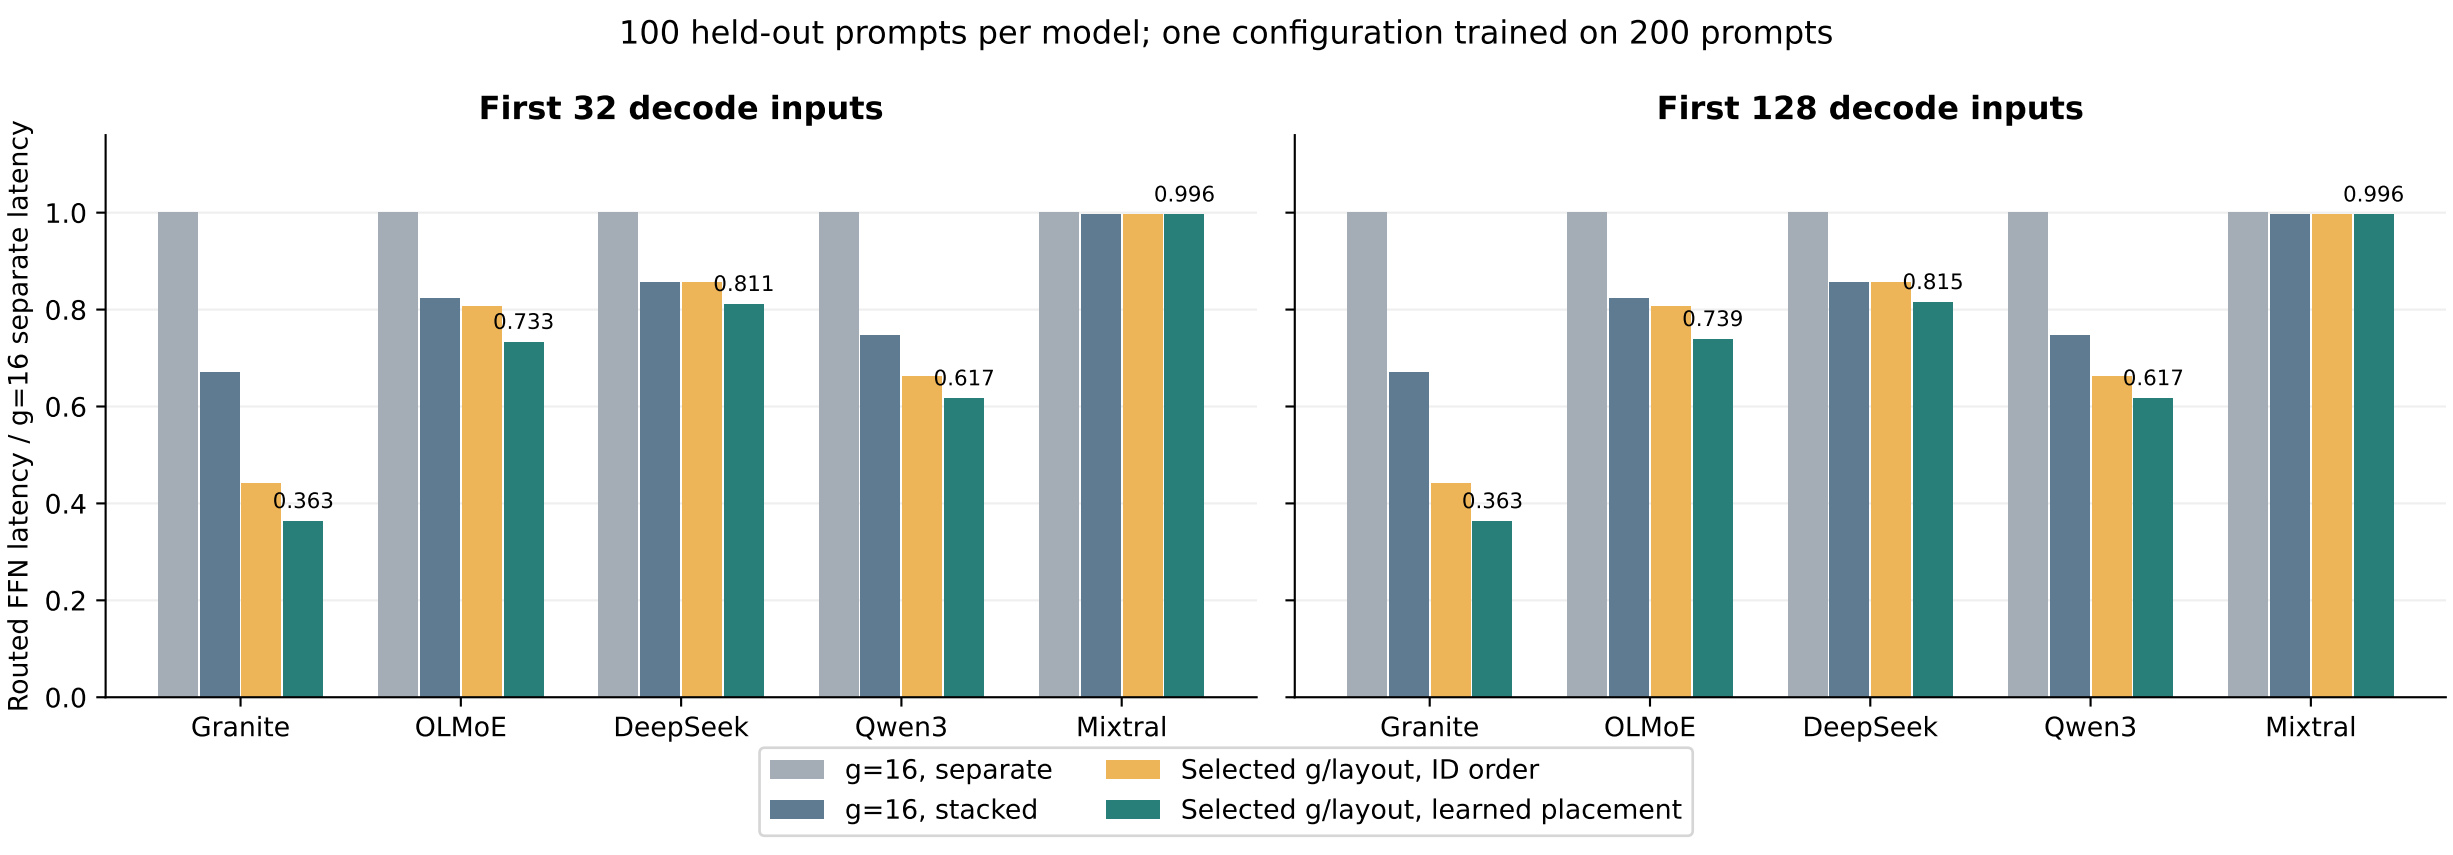
100 held-out prompts per model; one configuration trained on 200 prompts
First 32 decode inputs
First 128 decode inputs
Routed FFN latency / g=16 separate latency
0.0
0.2
0.4
0.6
0.8
1.0
0.363
0.733
0.811
0.617
0.996
0.363
0.739
0.815
0.617
0.996
Granite
OLMoE
DeepSeek
Qwen3
Mixtral
g=16, separate
g=16, stacked
Selected g/layout, ID order
Selected g/layout, learned placement