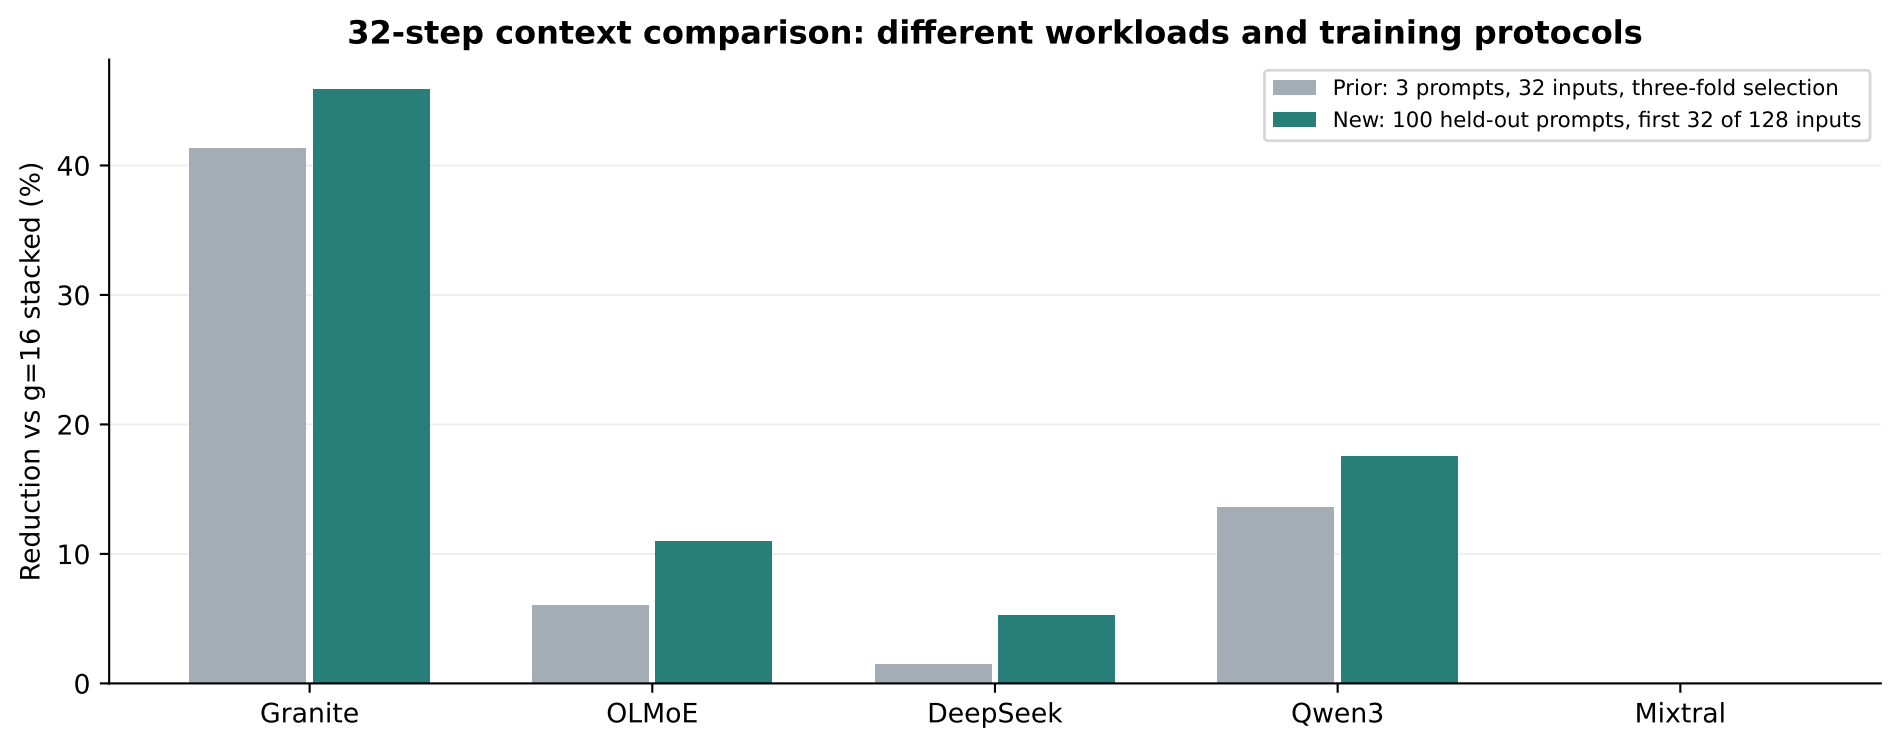
32-step context comparison: different workloads and training protocols
Prior: 3 prompts, 32 inputs, three-fold selection
New: 100 held-out prompts, first 32 of 128 inputs
Reduction vs g=16 stacked (%)
0
10
20
30
40
Granite
OLMoE
DeepSeek
Qwen3
Mixtral