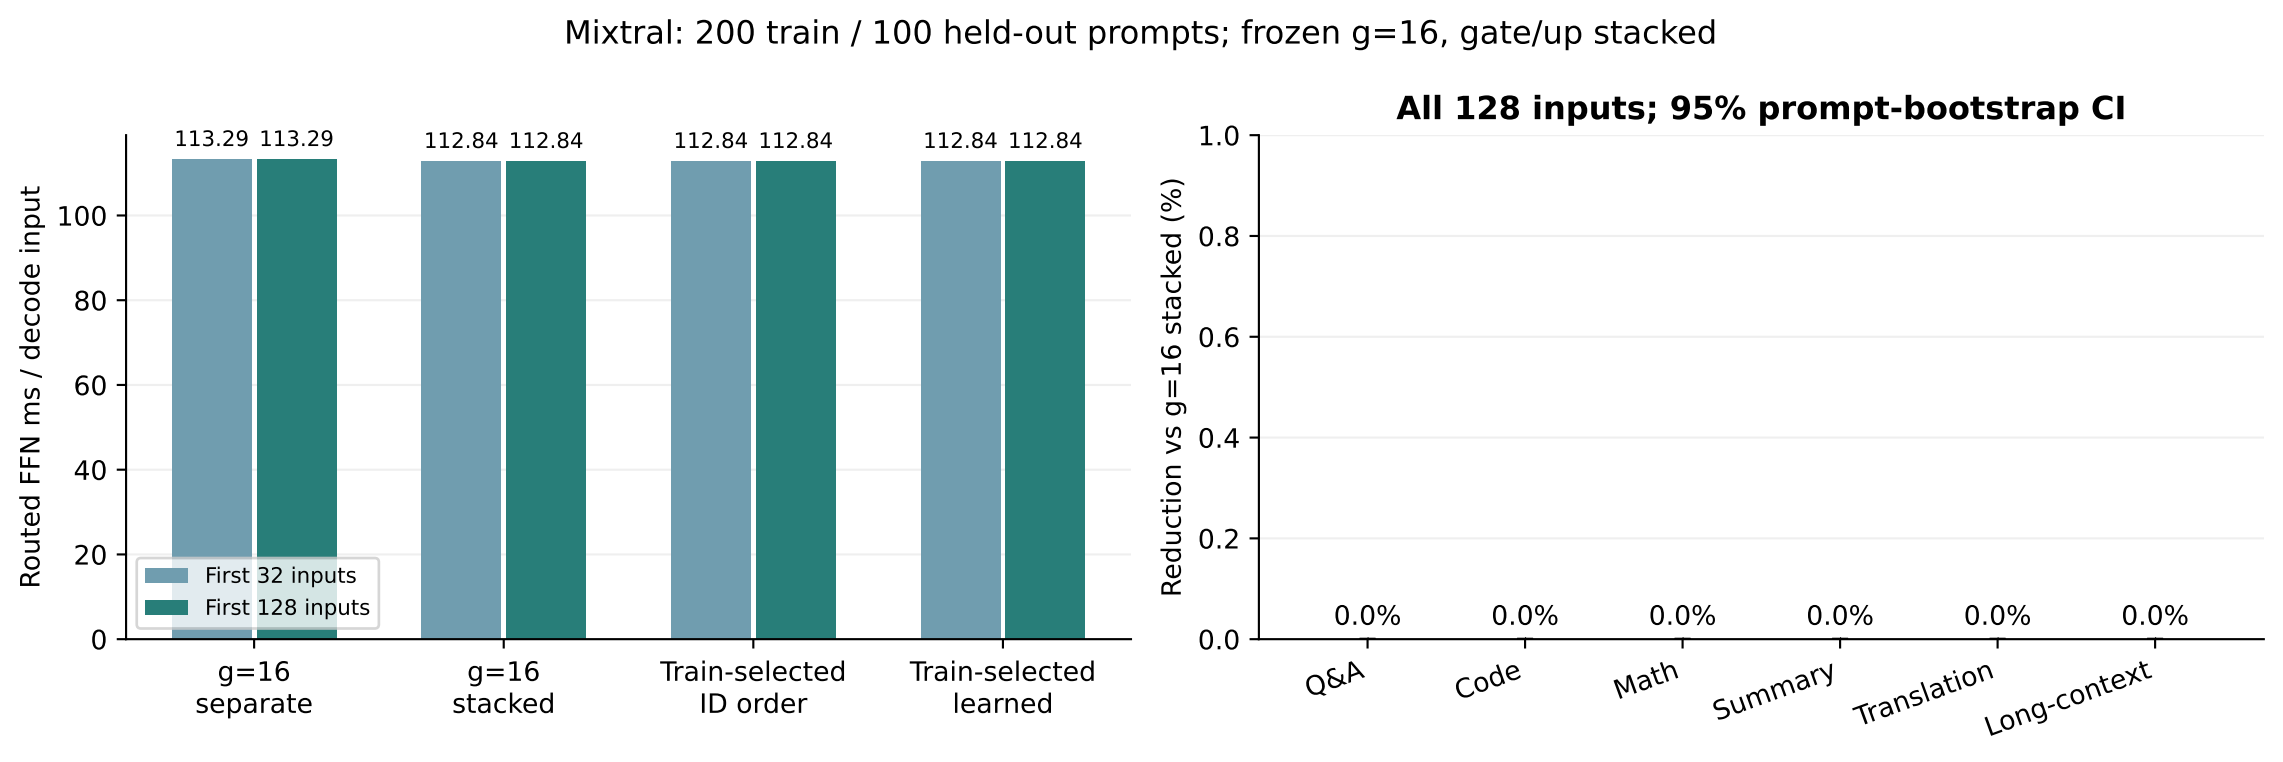
Mixtral: 200 train / 100 held-out prompts; frozen g=16, gate/up stacked
Routed FFN ms / decode input
113.29 113.29
112.84 112.84
112.84 112.84
112.84 112.84
First 32 inputs
First 128 inputs
g=16 separate
g=16 stacked
Train-selected ID order
Train-selected learned
All 128 inputs; 95% prompt-bootstrap CI
Reduction vs g=16 stacked (%)
0.0%
Q&A
Code
Math
Summary
Translation
Long-context