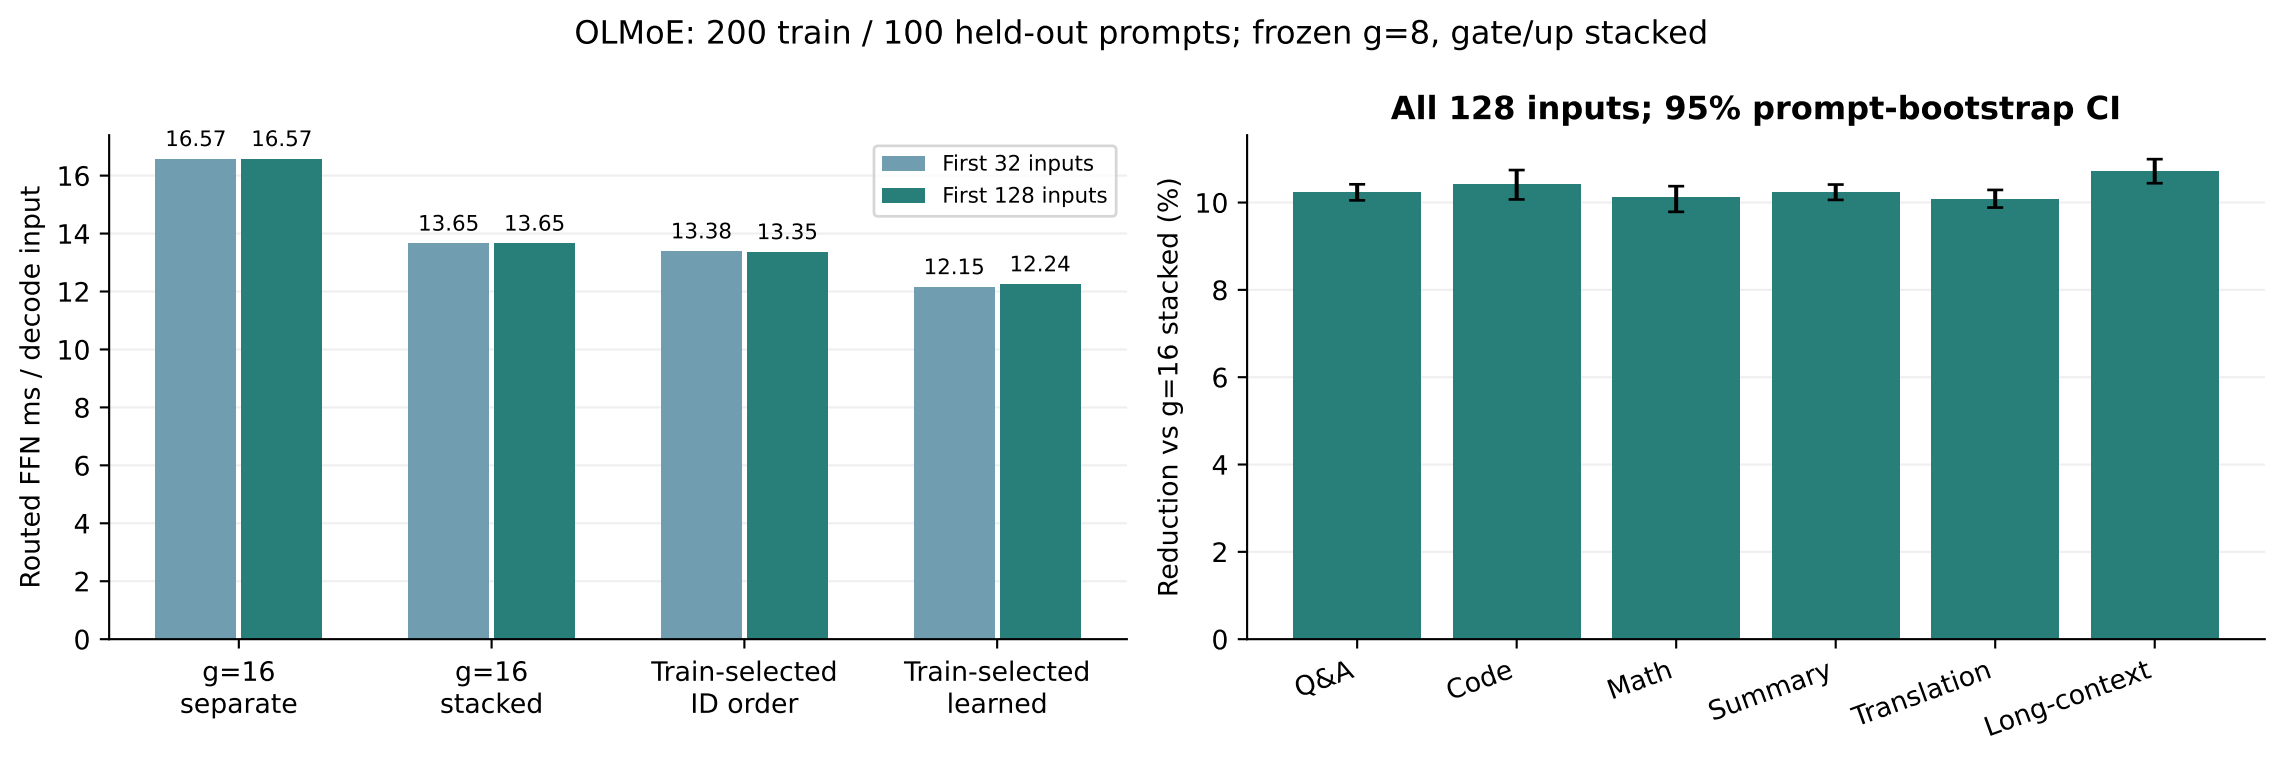
OLMoE: 200 train / 100 held-out prompts; frozen g=8, gate/up stacked
Routed FFN ms / decode input
First 32 inputs
First 128 inputs
16.57
16.57
13.65
13.65
13.38
13.35
12.15
12.24
g=16 separate
g=16 stacked
Train-selected ID order
Train-selected learned
All 128 inputs; 95% prompt-bootstrap CI
Reduction vs g=16 stacked (%)
Q&A
Code
Math
Summary
Translation
Long-context