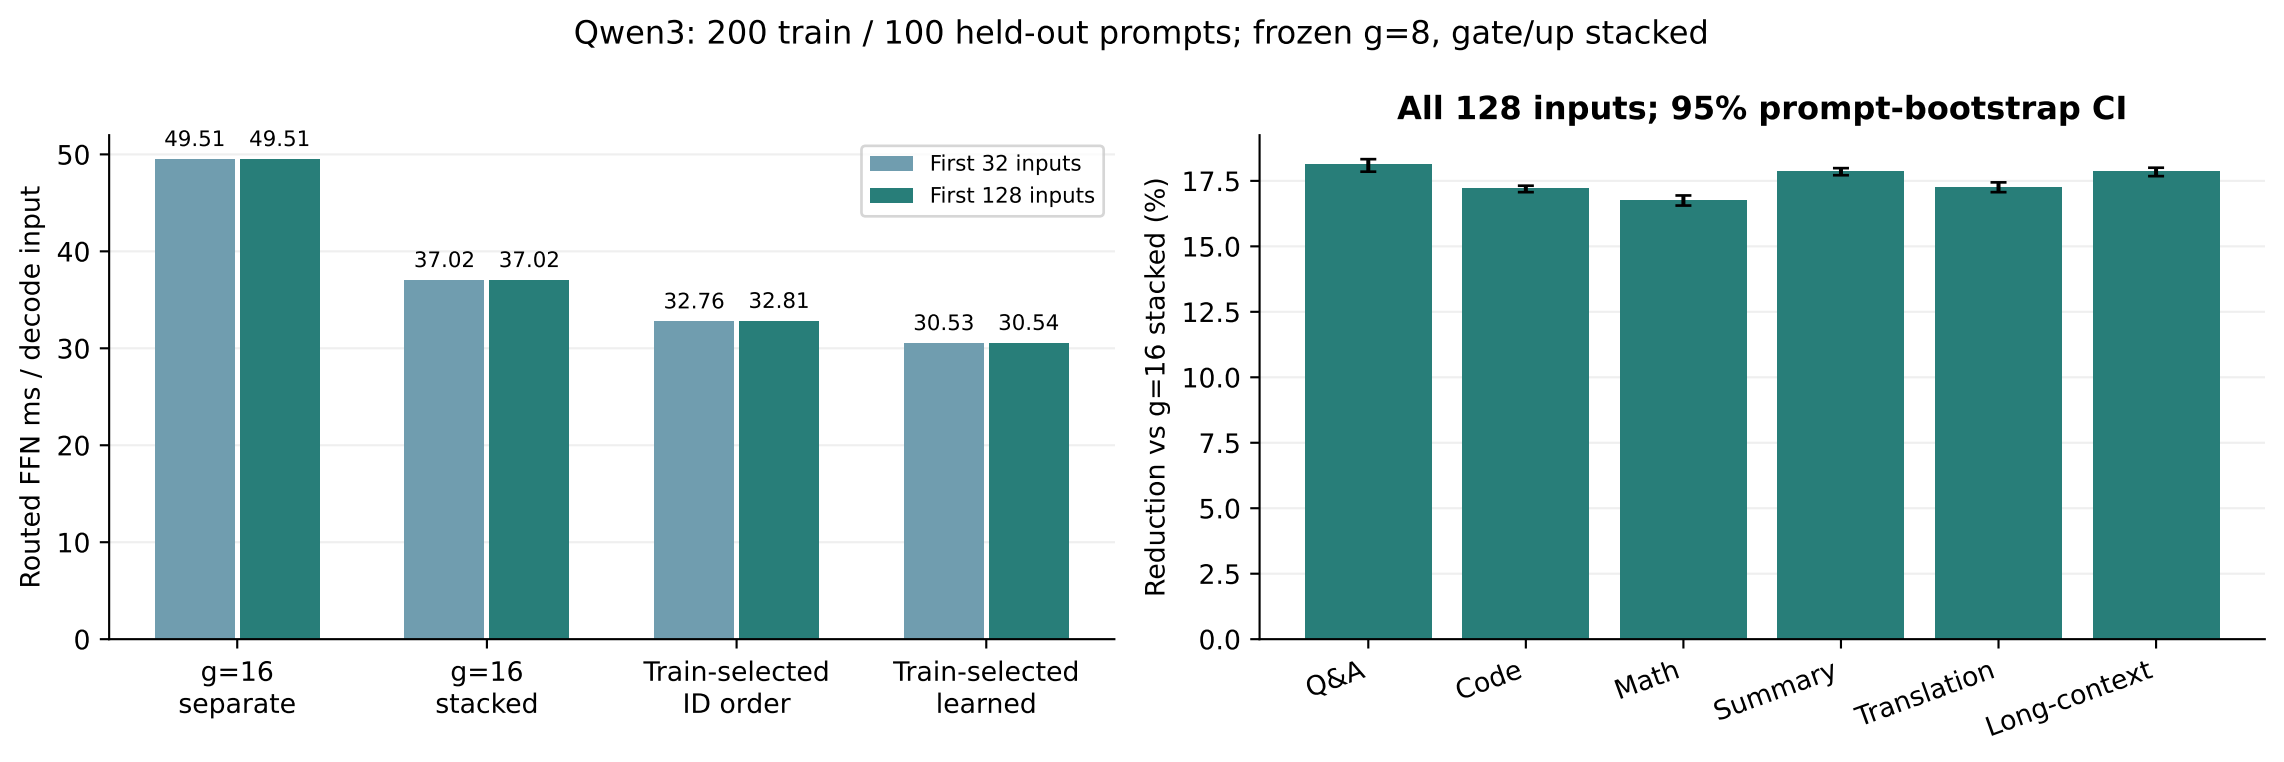
Qwen3: 200 train / 100 held-out prompts; frozen g=8, gate/up stacked
Routed FFN ms / decode input
First 32 inputs
First 128 inputs
49.51
49.51
37.02
37.02
32.76
32.81
30.53
30.54
g=16 separate
g=16 stacked
Train-selected ID order
Train-selected learned
All 128 inputs; 95% prompt-bootstrap CI
Reduction vs g=16 stacked (%)
Q&A
Code
Math
Summary
Translation
Long-context