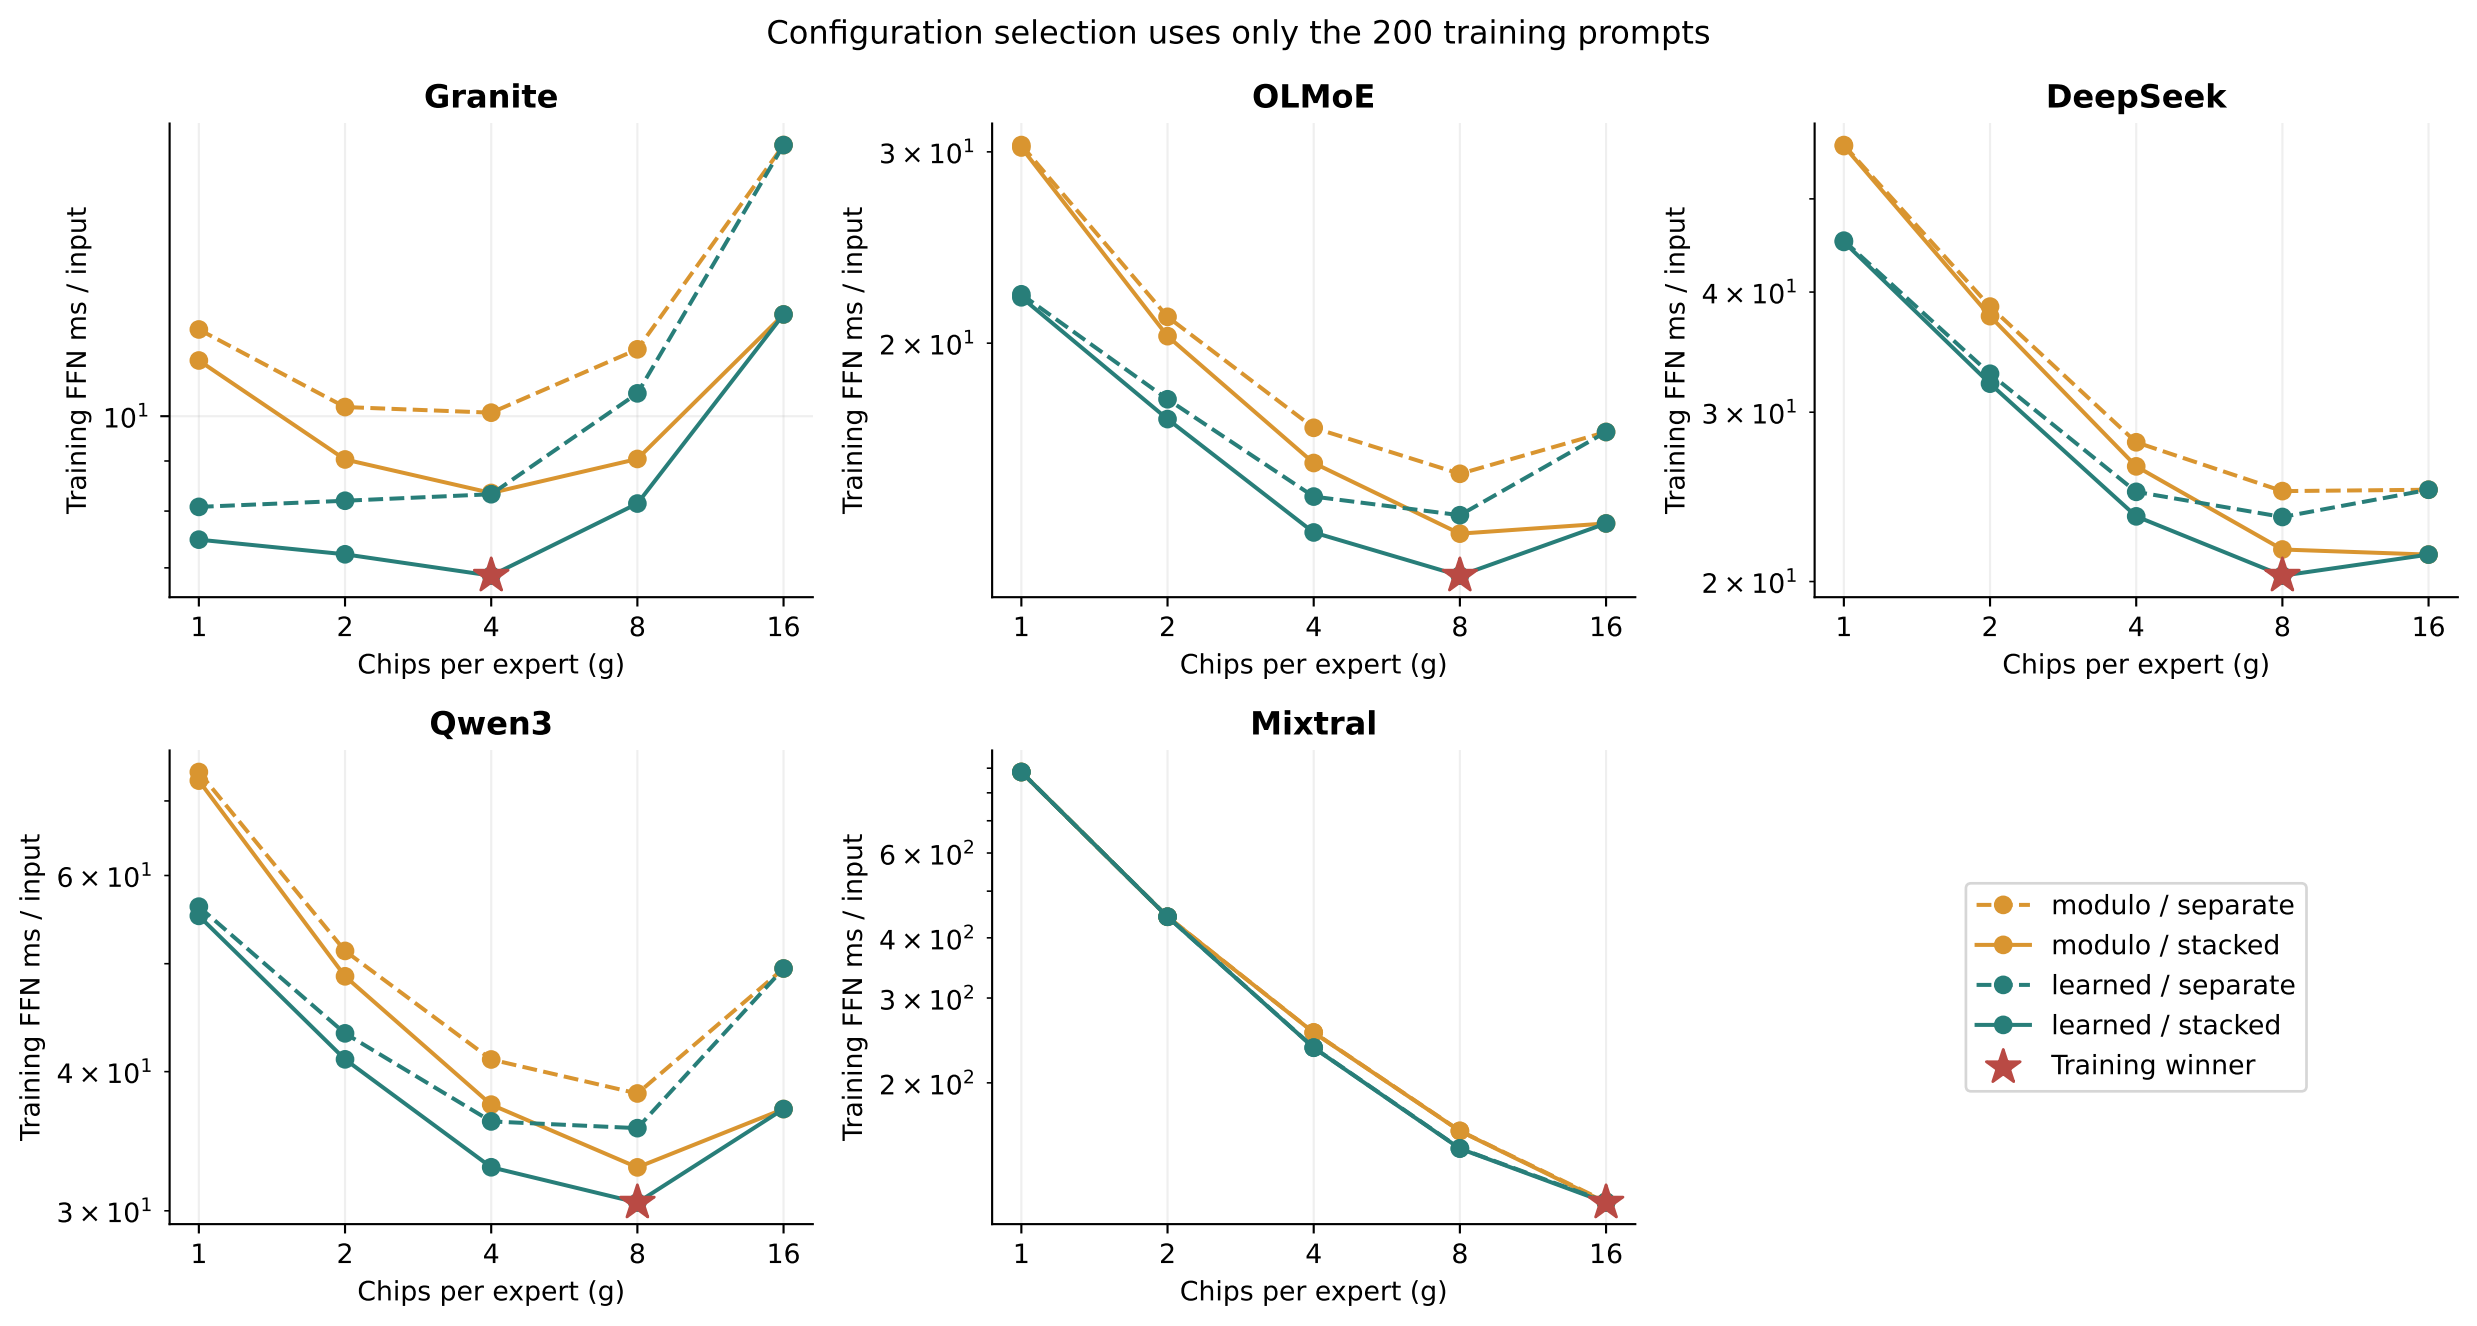
Configuration selection uses only the 200 training prompts
Granite
OLMoE
DeepSeek
Qwen3
Mixtral
Training FFN ms / input
Chips per expert (g)
1
2
4
8
16
10^1
2×10^1
3×10^1
4×10^1
6×10^1
2×10^2
3×10^2
4×10^2
6×10^2
modulo / separate
modulo / stacked
learned / separate
learned / stacked
Training winner