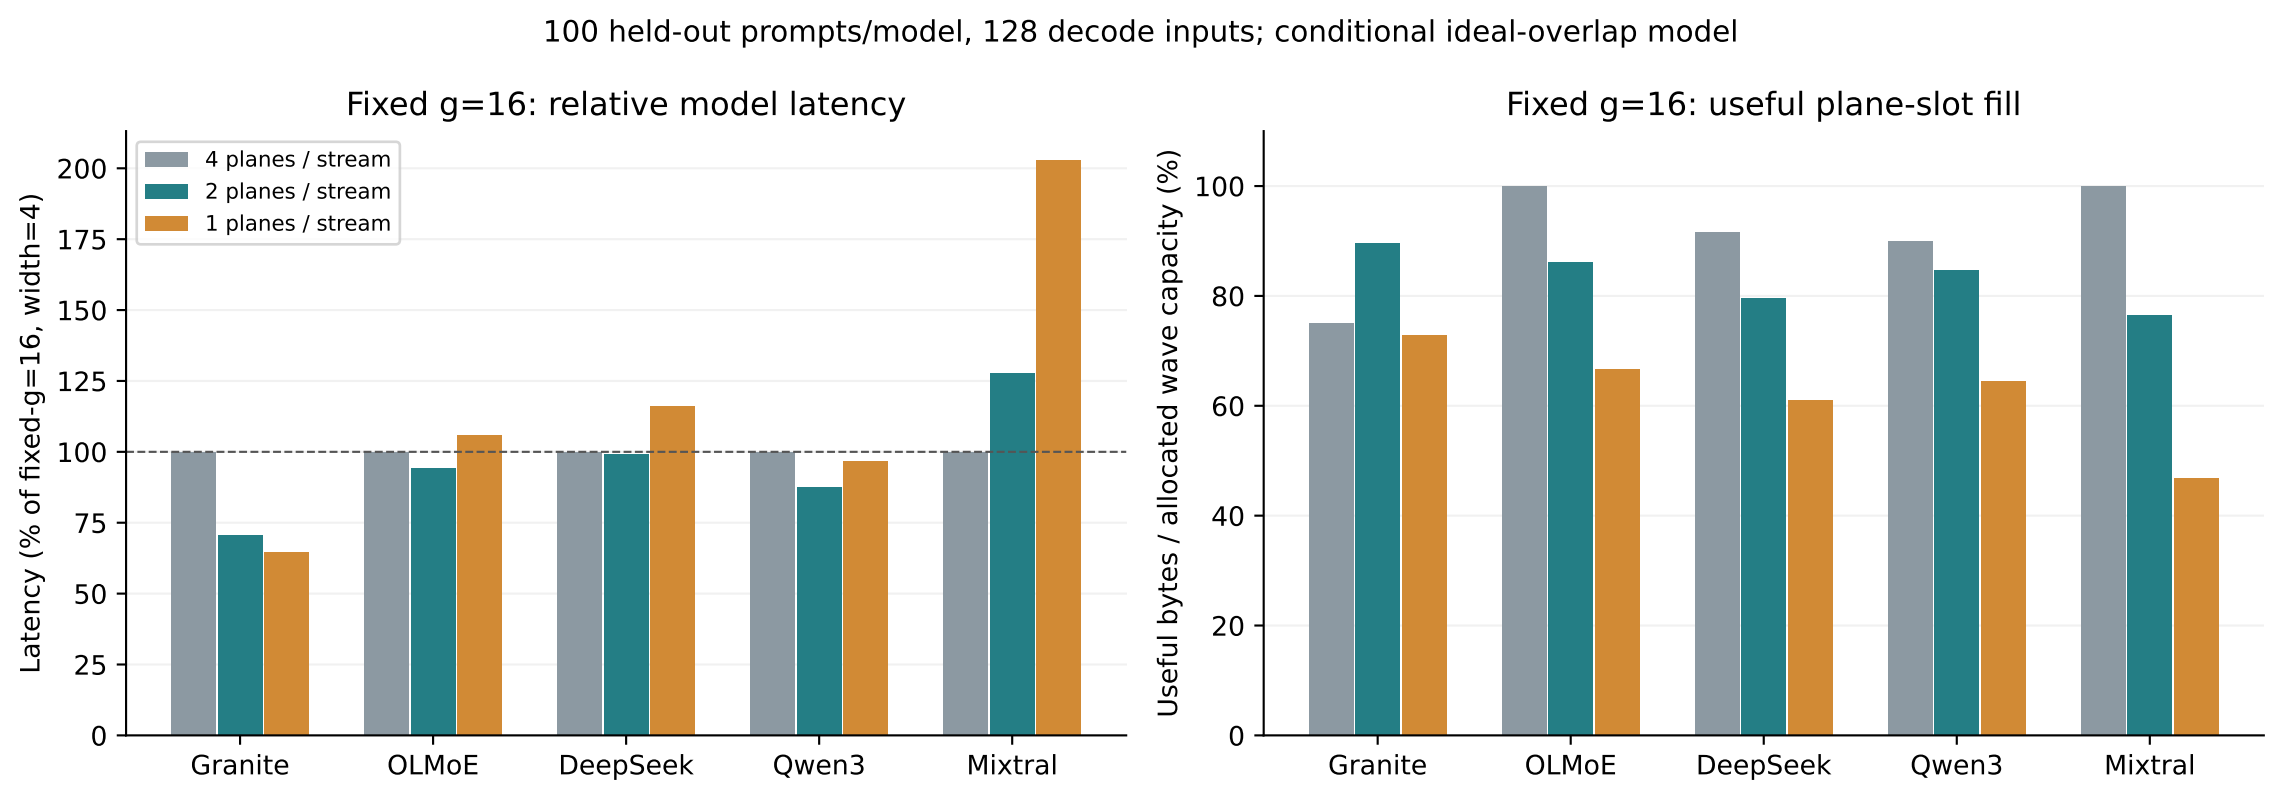
100 held-out prompts/model, 128 decode inputs; conditional ideal-overlap model
Fixed g=16: relative model latency
Latency (% of fixed-g=16, width=4)
4 planes / stream
2 planes / stream
1 planes / stream
0
25
50
75
100
125
150
175
200
Granite
OLMoE
DeepSeek
Qwen3
Mixtral
Fixed g=16: useful plane-slot fill
Useful bytes / allocated wave capacity (%)
0
20
40
60
80
100
Granite
OLMoE
DeepSeek
Qwen3
Mixtral