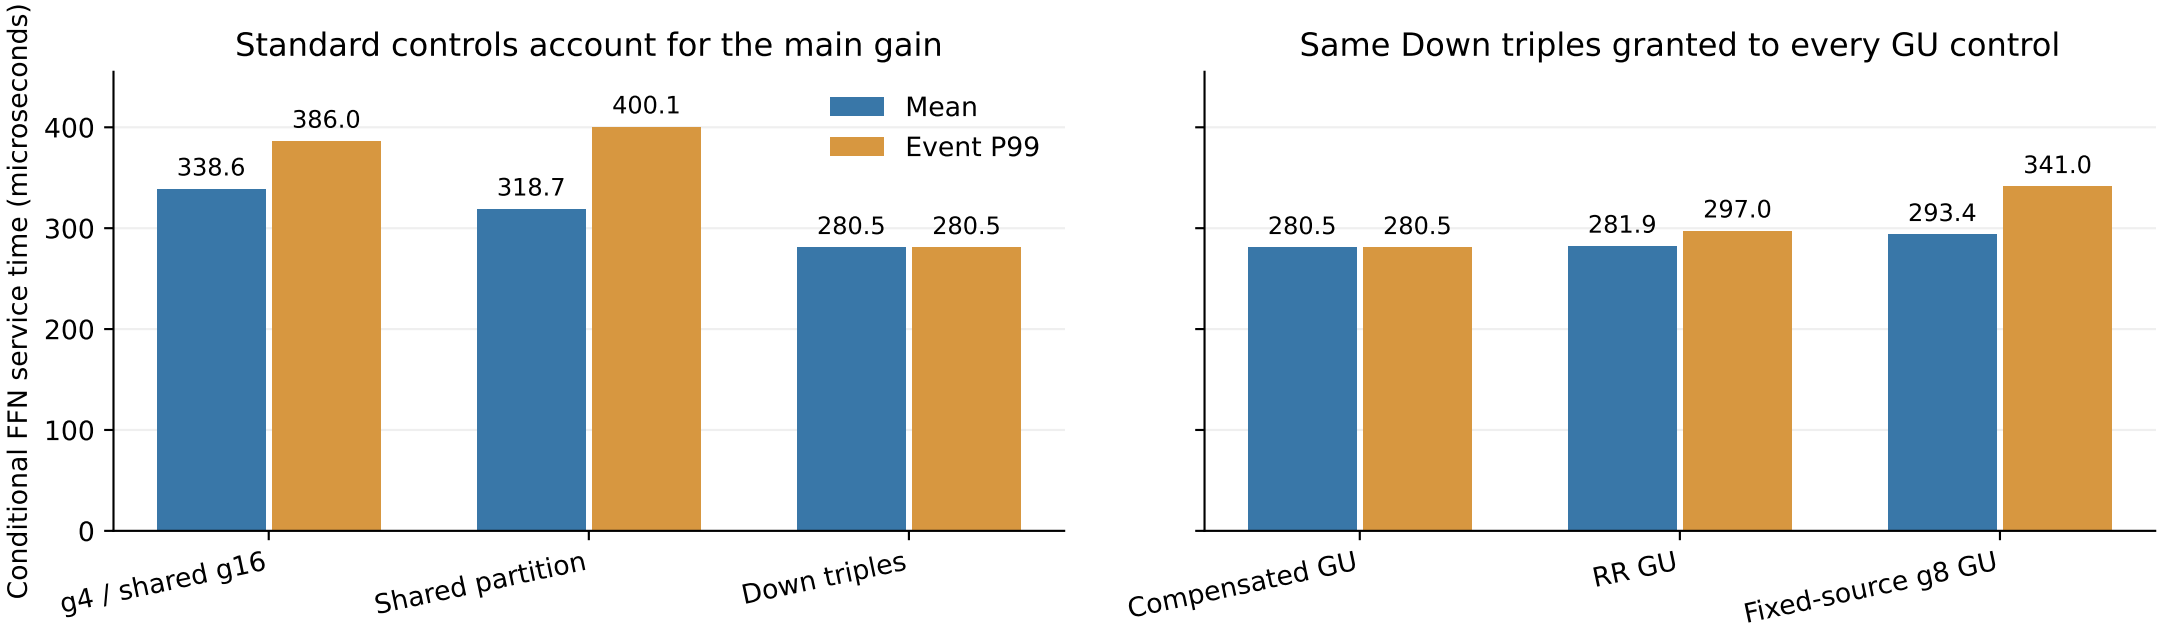

Down triples add 96 MiB/layer (148.5 to 244.5 MiB). Nominal copy500 / CR9.7us model; not measured device or full inference.
512 aligned inputs from 32 prompts; 1,536 layer-events. P99 is pooled layer-event latency, not a queueing SLO.